

Weekly Progress Report

1 – 8 September 2026 · Shishir Timalsina

Six workstreams were delivered this week: the Dashboard was rebuilt on real collected data, model access was put behind a single accounted seam and opened to administrator-supplied API keys, Mode Push scoring was rewritten with rarity weighting and given a results view, a production collection failure was diagnosed and fixed, and the interface was tightened across the shell and the digest. Four pull requests merged; one branch remains open.

10 commits	4 PRs merged	+5,050 lines added	16 new modules	+98 tests added	775 tests passing
----------------------	------------------------	------------------------------	--------------------------	---------------------------	-----------------------------

Figures are measured from the repository over the range 1–8 September, not estimated. Test total is for main; the open interface branch adds four more.

DELIVERED · PR #23

Dashboard rebuilt on collected data

The Dashboard tab was previously built entirely on invented numbers: a hardcoded twelve-week series, fabricated sector totals, a literal **20** for the watchlist count, and an “↑ trending” delta on every tile with nothing actually compared. It was the most confident-looking screen in the product and the only one where nothing on it was true — and the invented series ran an order of magnitude high, showing 847 Australian roles a week against a real 73.

It is now counted from the signals table. Three decisions shaped the replacement:

- **A point per collection, not per calendar week.** The pipeline runs weekly so the two usually coincide, but when a run is missed a calendar chart must draw something for the gap, and every option misleads — a zero says nobody hired, a joined line invents a measurement, and repeating the last value states it twice.
- **Movement is measured or absent.** Where there is no earlier collection to compare against, the tile says so rather than showing a direction it cannot justify.
- **Only classified signals count.** A collected but unread row has no sector or region, so including it would move the totals without being able to say where.

Fourteen tests cover the page, most of them about the two ways this kind of screen lies when data is thin: claiming a period it does not have, and showing a movement it has not measured.

DELIVERED · PR #24

Model access: one accounted seam, and keys an administrator can set

Usage accounting was reporting a small fraction of real spend. Only Mode Push went through the counting seam; classification, Market Pulse and Mode Publish each built their own client and were invisible to it. All four now route through `caller_for()`, so the figure on screen is the actual total, and those three become genuinely switchable to another provider rather than only appearing so.

Provider API keys can now be entered in the Admin panel. Keys were environment-only, so changing one meant editing a GitHub secret and waiting for a redeploy — the wrong shape for a credential that gets rotated precisely when it has leaked or run out of allowance.

- Keys are **encrypted before storage** with a secret held in the environment and never in the database, so a database dump yields ciphertext rather than working credentials.

- A deployment without that secret **refuses to store and says which variable to set**, rather than quietly writing plaintext an administrator would believe was protected.
- The key is **never returned**: the panel shows the last four characters, which identifies it to somebody already holding it and is useless to anybody else.
- **Test** spends one real call, because a stored key is not necessarily a working one — truncated by a paste, revoked, or a project with the API switched off all look identical until a run fails at five on a Monday morning.

DELIVERED · PR #25

Mode Push: rarity-weighted scoring, a results view, and outcome capture

The scorer counted every skill overlap equally, so a candidate and a company both saying “lean” was worth what both saying “HAZOP” was worth. Nearly every advert says lean; almost none say HAZOP. Skill matches are now weighted by inverse document frequency over the run’s own adverts — the standard approach in search and applicant tracking — and by what kind of skill it is, since a certification is a gate and a method is a turn of phrase.

Rarity is measured per run rather than from a fixed table, because what is unusual in Papua New Guinean mining is not what is unusual in Australian defence. Below 25 adverts it does not apply at all and the interface says so, since document frequency over a handful of postings is not a measurement.

Alongside it:

- A **skills taxonomy** with aliases, so “P6” and “Primavera P6” are one skill rather than a candidate scoring lower for writing it the other way. The previous list was 25 literal substrings that both missed and over-matched.
- **Job titles** now compare discipline with seniority words removed. “Senior Planner” against “Planner” was losing points on the title and again on seniority — the same fact charged twice.
- **A dedicated results view with a detail drawer**. Results previously filled a section beneath a long form, so reading them meant scrolling past the fields that produced them. Each company now opens to show every contributor as a bar, the contributors that could *not* be judged and why, and every claimed skill — matched or not — with its rarity.
- **Outcome capture**. Every weight in the model is judgement: nobody has been placed through this, so nothing is calibrated. Each ranked company can now be marked contacted, placed or not relevant, storing the score as *it stood at that moment*.

On accuracy: no accuracy figure is reported, because none can honestly be computed yet. The summary deliberately returns counts and a threshold rather than a success rate — a precision figure over four decisions is not a rough measurement, it is a wrong one. Freezing the score with the decision is what makes later calibration meaningful: recomputing it would judge a decision against a model that did not exist when it was taken.

DIAGNOSED AND FIXED · PR #26

Production incident: no PNG collection, and a missing Market Pulse

Monday’s scheduled run reported healthy while collecting **zero** Papua New Guinean job signals and archiving a digest with no Market Pulse. Both were traced to root cause from the production database rather than guessed at.

Empty is not a value. The deployment writes the server’s environment file from GitHub variables, and an unset variable expands to nothing — producing `PNGWORKFORCE_BASE_URL=`, present and empty. Python’s `os.environ.get(name, default)` applies the default only when the name is *missing*, so the empty string overrode a correct default. The scraper had no URL, returned an empty list, and the run reported success.

The commit that introduced that passthrough was itself the cause: it added the code default *and* the line that overrides it, and was verified with the variable removed from the environment — the one state the same change made impossible in production. It worked on every developer machine, where the variable is genuinely absent.

The Market Pulse was generated but never archived. The digest payload is built before the pulse is generated, so the archived copy always recorded it as absent. It only ever looked correct when a window happened to be

processed twice — which is why every digest examined by hand had appeared fine.

IN REVIEW · BRANCH OPEN

Interface pass across the shell, the digest and the admin screens

- The sign-in page now describes **Mode Publish** alongside Monitor and Push — the third mode was invisible on the one screen every user sees first.
- The placeholder logo was removed. No mark has been chosen, and a placeholder shown long enough stops reading as one.
- Admin navigation renamed for what each page manages. **“Tokens & cost” was wrong twice over**: the page counts requests against a daily allowance and never tokens, and it has since become where API keys are entered and a model chosen per job.
- **Market Pulse** moved onto the same card chassis as the rest of the page; it had its own border weight, header rule and gradient, which made the only written passage on the screen read as an insert from another product.
- The week selector was rebuilt to match the application’s own controls, and the “Archived” and “Live data” chips removed — both said the same thing on every visit. **“Sample data” was deliberately kept**: it is the only thing distinguishing a real week from a demonstration one.
- A **switched-off source no longer reports “Collecting”**. Status was derived from the age of the newest signal alone, so SEEK showed a green chip beside a toggle reading Off.
- Long explanatory prose on the admin screens folds into disclosures titled for the question each answers.

QUALITY

Defects found and corrected

Several of these were found by tests written against the change rather than by using the product, and three were self-inflicted regressions caught before release.

Defect	Consequence had it shipped
Rarity weighting applied to both sides of the ratio	Cancelled out entirely — one rare skill matched would have scored exactly as one common skill matched, which is the flat model it replaced.
Evaluation harness cleared a counter key that no longer existed	Each evaluation run inherited the previous run's quota count and could refuse to classify, reporting an exhausted allowance that was fine.
Mode Publish still gated on the Gemini environment key	Would have refused a perfectly valid key entered in the Admin panel, naming a variable that had stopped being the authority.
Deployment secret never reached the container	Setting it would have appeared to work and changed nothing; the panel would have refused every key with no indication why.
Verification endpoint counted only successful calls	The Test button could have silently exhausted a free tier — the exact hole the counting seam was built to close.
Time-dependent test failing near midnight UTC	Intermittent CI failure, reproduced by simulation across the clock and fixed by anchoring the fixture.
Usage row collapsed its grid tracks to zero width	Text painted outside its own box and was one line-length from wrapping into a column narrower than a word.
Status marker style never defined	Every not-configured or retired source rendered a blank where other rows carry a marker.

ASSURANCE

Testing and verification

Ninety-eight net new tests were added, taking the suite to **775 passing** with 10 skipped. Coverage was written to pin behaviour rather than implementation, and in two cases each fix was confirmed by removing it and observing the test fail — which mattered here, because one of the bugs fixed this week had originally shipped *with* a passing verification.

- A guard asserting that **no module outside the seam imports a provider SDK**, which is the check that would have caught the three uncounted callers in the first place.
- A contract test pinning the payload the scoring drawer reads **by field name** — the API is Python and the interface TypeScript, so a rename type-checks cleanly on both sides and surfaces as an empty panel.
- A test asserting the lazily-created database table matches the one in the schema, so two definitions of one table cannot drift.
- Interface changes were verified in a browser against the real stylesheet, with geometry measured rather than eyeballed — which is how two of the layout defects above were found, both of which looked correct in a screenshot.

OUTLOOK

Open items and next week

- **Merge the interface branch** and confirm the digest and admin screens in the deployed application.
- **Confirm the PNGworkforce fix on Monday's scheduled run.** The mechanism is proven and the fix is in, but the container log for the failed run was never read, so the diagnosis remains inferred rather than observed.
- **Scrape concurrency.** The four sources are collected sequentially, so a run takes the sum of their times rather than the longest. Two share global crawler state, so this needs care rather than a single change.
- **Begin accumulating Mode Push outcomes**, which is the only route to a scoring model calibrated against placements rather than judgement.

Prepared from the repository history for the period 1–8 September 2026. Commit range, diffstat, test counts and defect list are measured; no figure in this document is estimated.